

GROUP-4

CAPSTONE PROJECT-1

Customer Segmentation

- RFM Analysis
- K-Means Modelling
- Cohort Analysis

Contents

- Group Members
- Main Steps
- Meeting Schedule
- Methodology
 - RFM Analysis
 - K-Means Modelling
 - Cohort Analysis



Group Members

- C8124-Jack Sellow
- C8125-Mustafa
- C8250-Onur
- C8278-Engin
- C8292-Ken
- C8301-Sam
- C8307-Hakan
- C8315-Halit
- C8399-Benjamin
- C8492-Semih
- C9231-Hüseyin

Main Steps

1. Data Cleaning & EDA
2. RFM Analysis
3. RFM Segmentation
4. K-Means Modelling
5. Cohort Analysis

Schedule

04-06 December 2021 :	Individual study
06 December 2021 : 13:00 (IST)	Data Cleaning & EDA
07 December 2021 : 20:00 (IST)	RFM Analysis & Segmentation
08 December 2021 : 20:00 (IST)	K-Means Modelling & Cohort Analysis
09 December 2021 : 13:00 (IST)	Review

Methodology

Analysis

Columns : 8

Rows : 541.909

Duplicated: 5.268

EDA

RFM Analysis

K-Means

Cohort Analysis

#	Column	Non-Null	Count	Dtype
0	InvoiceNo	541909	non-null	object
1	StockCode	541909	non-null	object
2	Description	540455	non-null	object
3	Quantity	541909	non-null	int64
4	InvoiceDate	541909	non-null	datetime64[ns]
5	UnitPrice	541909	non-null	float64
6	CustomerID	406829	non-null	float64
7	Country	541909	non-null	object

Missing Values:

	Missing_Number	Missing_Percent
CustomerID	135080	0.249267
Description	1454	0.002683

Number of Uniques:

InvoiceNo	25900
StockCode	4070
Description	4223
Quantity	722
InvoiceDate	23260
UnitPrice	1630
CustomerID	4372
Country	38

Descriptive Statistics

	Quantity	UnitPrice	CustomerID
count	541909.00	541909.00	406829.00
mean	9.55	4.61	15287.69
std	218.08	96.76	1713.60
min	-80995.00	-11062.06	12346.00
25%	1.00	1.25	13953.00
50%	3.00	2.08	15152.00
75%	10.00	4.13	16791.00
max	80995.00	38970.00	18287.00

Descriptive Statistics (Categorical Columns)

	count	unique	top	freq
InvoiceNo	541909	25900	573585	1114
StockCode	541909	4070	85123A	2313
Description	540455	4223	WHITE HANGING HEART T-LIGHT HOLDER	2369
Country	541909	38	United Kingdom	495478

Methodology

Analysis

1. InvoiceNo
2. StockCode
3. Description
4. Quantity
5. InvoiceDate
6. UnitPrice
7. CustomerID
8. Country

25.900 unique invoices

3.836 cancelled invoices

22.064 others

```
df["invoiceno"].str.startswith('C').value_counts(normalize = True)*100
```

False	98.286059
True	1.713941

Methodology

Analysis

1. InvoiceNo

AMAZONFEE : (32 of the 34 AMAZONFEE invoices have been cancelled)

2. StockCode

BANK CHARGES : (25 of the 37 BANK CHARGE invoices have been cancelled)

3. Description

D : Discount (All of the 77 D invoices have been cancelled)

4. Quantity

M / m : Manual (244 of the 572 M invoices have been cancelled)

5. InvoiceDate

S : Samples (61 of the 63 S invoices have been cancelled. InvNo : 549684 – 572849)

6. UnitPrice

B : Adjust bad debt (2 of the 3 B invoices have been cancelled)

7. CustomerID

C2 : Carriage (2 of the 244 C2 invoices have been cancelled)

8. Country

DOR : DOTCOM POSTAGE (1 of the 710 DOT invoices has been cancelled)

POST : POSTAGE (126 of the 1256 POST invoices have been cancelled)

PADS : PADS TO MATCH ALL CUSHIONS (Any of the 4 PADS invoices has not been cancelled)

CRUK : CRUK Commission (All of the 16 CRUK invoices have been cancelled)

Methodology

Analysis

1. InvoiceNo

2. StockCode

3. Description

4. Quantity

5. InvoiceDate

6. UnitPrice

7. CustomerID

8. Country

```
(df["Quantity"] < 0).value_counts()
```

False	526054
True	10587

A negative quantity indicates canceled invoices or returned items.

```
df[(df["InvoiceNo"].str.find("C").isnull()) & (df["Quantity"] < 0)][["Description"].unique()
```

```
array([nan, '?', 'check', 'damages', 'faulty', 'Dotcom sales',  
      'reverse 21/5/10 adjustment', 'mouldy, thrown away.', 'counted',  
      'Given away', 'Dotcom', 'label mix up', 'samples/damages',  
      'thrown away', 'incorrectly made-thrown away.', 'showroom', 'MIA',  
      'Dotcom set', 'wrongly sold as sets', 'Amazon sold sets',  
      'dotcom sold sets', 'wrongly sold sets', '? sold as sets?',  
      '?sold as sets?', 'Thrown away.', 'damages/display',  
      'damaged stock', 'broken', 'throw away', 'wrong barcode (22467)',  
      'wrong barcode', 'barcode problem', '?lost',  
      "thrown away-can't sell.", "thrown away-can't sell", 'damages?',  
      're dotcom quick fix.', "Dotcom sold in 6's", 'sold in set?',  
      'cracked', 'sold as 22467', 'Damaged',  
      'mystery! Only ever imported 1800',  
      'MERCHANT CHANDLER CREDIT ERROR, STO', 'POSSIBLE DAMAGES OR LOST?',  
      'damaged', 'DAMAGED', 'Display', 'Missing', 'wrong code?',  
      'wrong code', 'adjust', 'crushed', 'damages/showroom etc',  
      'samples', 'damages/credits from ASOS.',  
      'Not rcvd in 10/11/2010 delivery', 'Thrown away-rusty',
```

Methodology

Analysis

1. InvoiceNo

2. StockCode

3. Description

4. Quantity

5. InvoiceDate

6. UnitPrice

7. CustomerID

8. Country

```
print("Max Date : ", df["InvoiceDate"].max())
print("Min Date : ", df["InvoiceDate"].min())
```

```
Max Date : 2011-12-09 12:50:00
Min Date : 2010-12-01 08:26:00
```

```
df[df["UnitPrice"] < 0]
```

	InvoiceNo	StockCode	Description	Quantity	InvoiceDate	UnitPrice	CustomerID	Country
299983	A563186	B	Adjust bad debt	1	2011-08-12 14:51:00	-11062.06	NaN	United Kingdom
299984	A563187	B	Adjust bad debt	1	2011-08-12 14:52:00	-11062.06	NaN	United Kingdom

```
df[df["StockCode"] == "B"]
```

	InvoiceNo	StockCode	Description	Quantity	InvoiceDate	UnitPrice	CustomerID	Country
299982	A563185	B	Adjust bad debt	1	2011-08-12 14:50:00	11062.06	NaN	United Kingdom
299983	A563186	B	Adjust bad debt	1	2011-08-12 14:51:00	-11062.06	NaN	United Kingdom
299984	A563187	B	Adjust bad debt	1	2011-08-12 14:52:00	-11062.06	NaN	United Kingdom

Methodology

Analysis

1. InvoiceNo

2. StockCode

3. Description

4. Quantity

5. InvoiceDate

6. UnitPrice

7. CustomerID

8. Country

```
df[df.InvoiceNo.str.contains("C").isnull()].groupby(by="CustomerID").count()["InvoiceNo"].sort_values(ascending=False).head(3)
```

```
CustomerID
17841.0    7676
14911.0    5672
14096.0    5111
Name: InvoiceNo, dtype: int64
```

- **Customer 17841** is the most purchasing customer.

```
df.groupby(by="CustomerID").sum()["TotalCost"].sort_values(ascending=False).head(3)
```

```
CustomerID
14646.0    279489.02
18102.0    256438.49
17450.0    187322.17
Name: TotalCost, dtype: float64
```

- **Customer 14646** is the most revenue customer.

```
print("Number of average un-cancelled invoices per customer :")
print(df[df.InvoiceNo.str.contains("C").isnull()][df.InvoiceNo].unique().shape[0] / \
df[df.InvoiceNo.str.contains("C").isnull()][df.CustomerID].unique().shape[0])
```

```
Number of average un-cancelled invoices per customer :
5.083179723502304
```

Methodology

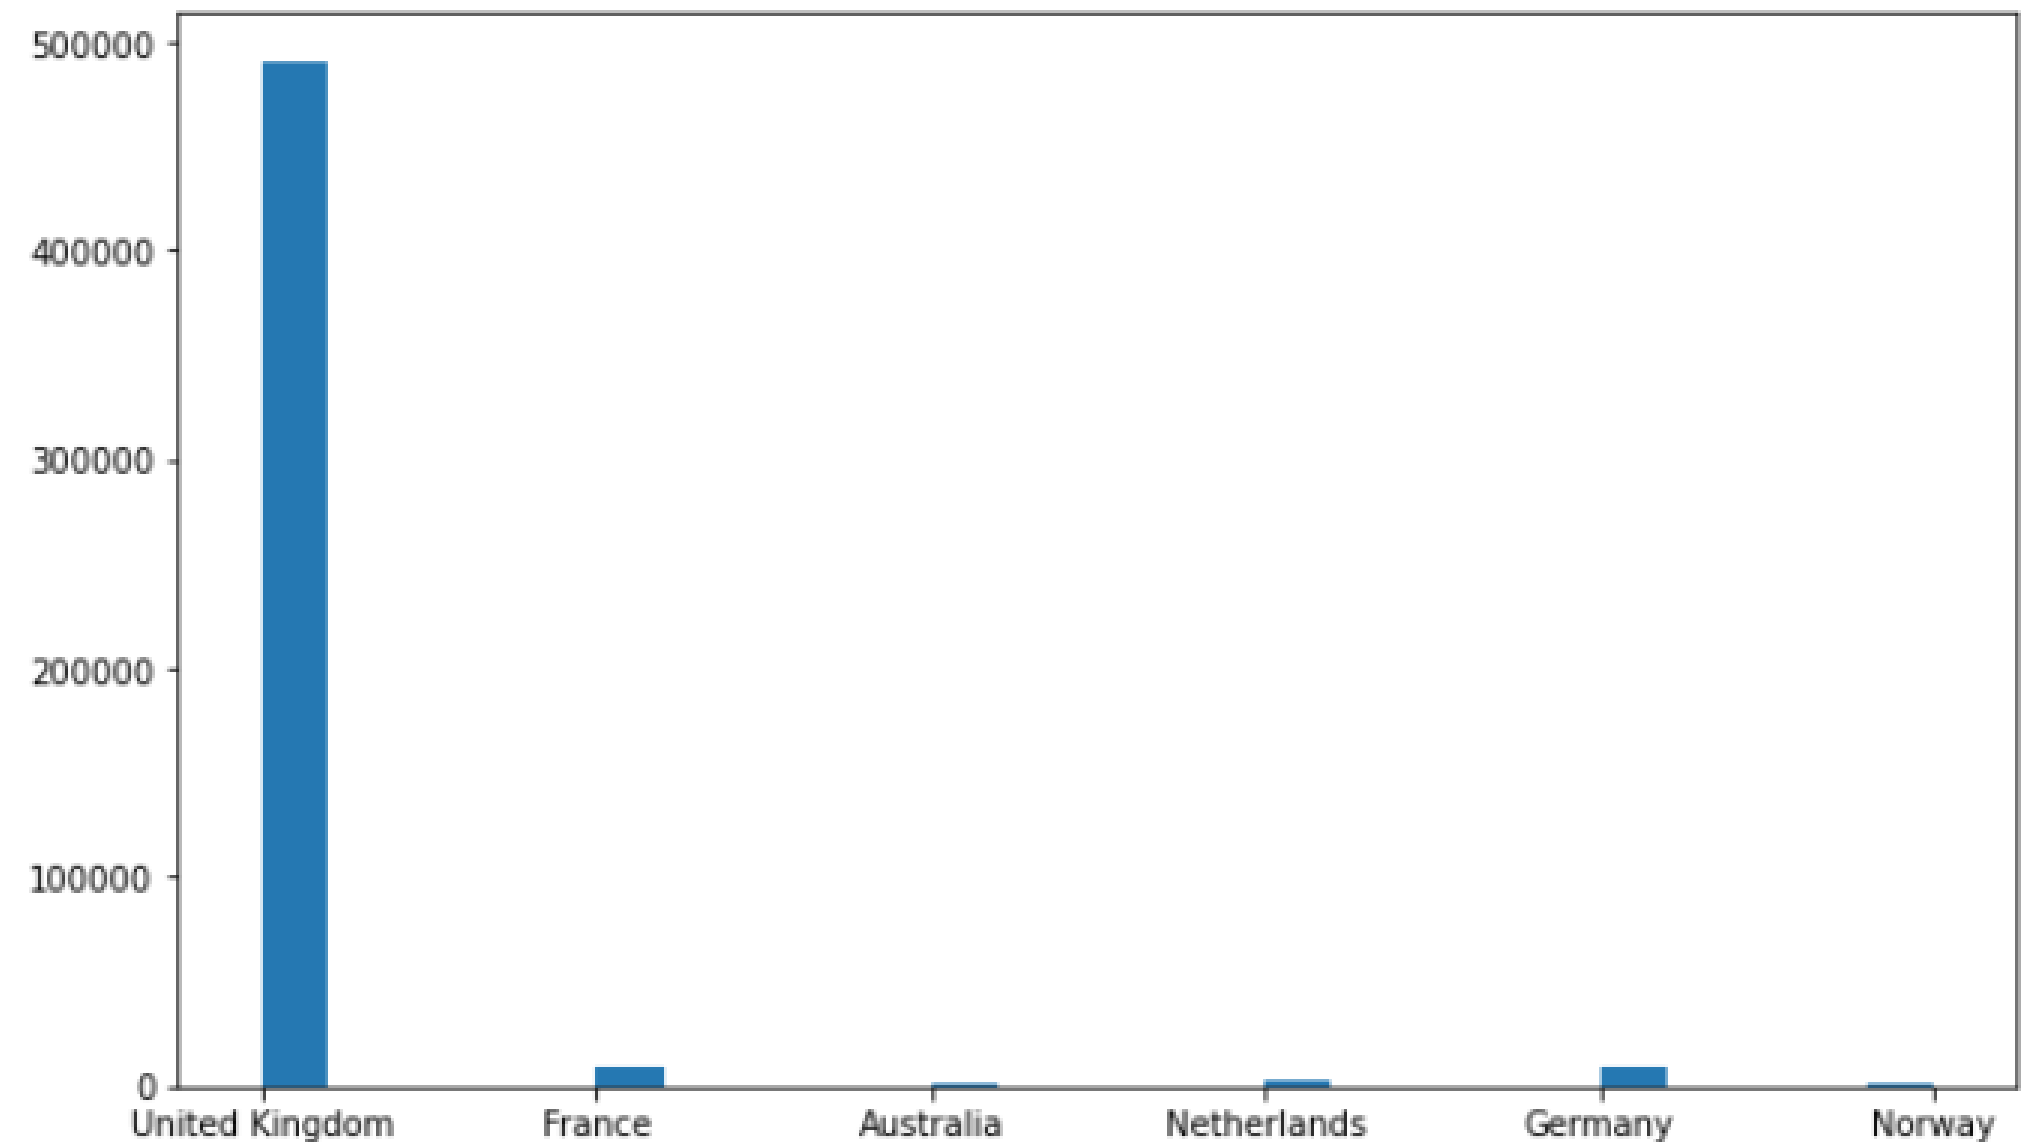
Analysis

1. InvoiceNo
2. StockCode
3. Description
4. Quantity
5. InvoiceDate
6. UnitPrice
7. CustomerID
8. Country

```
df["Country"].value_counts()
```

United Kingdom	490300
Germany	9480
France	8541
EIRE	8184
Spain	2528
Netherlands	2371
Belgium	2069
Switzerland	1994
Portugal	1510
Australia	1258
Norway	1086
Italy	803
Channel Islands	757
Finland	695
Cyprus	611
Sweden	461
Unspecified	442
Austria	401
Denmark	389
Japan	358
Poland	341
Israel	294
USA	291
Hong Kong	284
Singapore	229
Iceland	182
Canada	151
Greece	146
Malta	127

```
plt.hist(data=df, x="Country", bins=25, range=(0,5) );
```



Methodology

Analysis

EDA

RFM Analysis

K-Means

Cohort Analysis

Dropped

- Duplicates
- Cancelled Invoices
- Negative UnitPrices
- Negative Quantities
- Missing CustomerIDs

Rows

541.909



392.692

Methodology

Analysis

EDA

RFM Analysis

K-Means

Cohort Analysis

Total revenue
per country

```
df.groupby("country")['total_price'].sum().sort_values(ascending=False)
```

country	
United Kingdom	7285024.644
Netherlands	285446.340
EIRE	265262.460
Germany	228678.400
France	208934.310
Australia	138453.810
Spain	61558.560
Switzerland	56443.950
Belgium	41196.340

Total customer
per country

```
df.groupby("country")['customerid'].nunique().sort_values(ascending=False)
```

country	
United Kingdom	3920
Germany	94
France	87
Spain	30
Belgium	25
Switzerland	21
Portugal	19
Italy	14
Finland	12

Methodology

Analysis

EDA

RFM Analysis

K-Means

Cohort Analysis



Focused on the UK market

```
df_uk = df[df["country"]=="United Kingdom"]  
df_uk.head(1)
```

```
df_uk.shape  
  
(349203, 12)
```

2. What are the most popular products that are bought in the UK?

```
df_uk.groupby(["stockcode"])["quantity"].sum().sort_values(ascending=False).head(5)
```


stockcode	
23843	80995
23166	76919
84077	49086
22197	45609
85099B	41878

Methodology

Recency

Frequency

Monetary



RFM Metrics



RECENCY

The freshness of the customer activity, be it purchases or visits

E.g. Time since last order or last engaged with the product



FREQUENCY

The frequency of the customer transactions or visits

E.g. Total number of transactions or average time between transactions/engaged visits



MONETARY

The intention of customer to spend or purchasing power of customer

E.g. Total or average transactions value

CUSTOMER ID	RECENCY (DAY)	FREQUENCY (NUMBER)	MONETARY (TOTAL)
1	4	6	540
2	6	11	940
3	46	1	35
4	23	3	65
5	15	4	179
6	32	2	56
7	7	3	140
8	50	1	950
9	34	15	2630
10	10	5	191
11	3	8	845
12	1	10	1510
13	27	3	54
14	18	2	40
15	5	1	25

Analysis

EDA

RFM Analysis

K-Means

Cohort Analysis

Methodology

Analysis

EDA

RFM Analysis

K-Means

Cohort Analysis

Recency

ref_date = 2011-12-09

```
df_uk["customer_recency"] = df_uk["ref_date"] - df_uk["last_purchase_date"]
```

```
customer_recency = pd.DataFrame(df_uk.groupby('customerid')['recency_value'].min())
```

Frequency

```
customer_frequency = pd.DataFrame(df_uk.groupby('customerid')['invoiceno'].nunique())
```

```
df_uk['customer_frequency'] = df_uk.groupby('customerid')['invoiceno'].transform('count')
```

Monetary

```
customer_monetary = pd.DataFrame(df_uk.groupby('customerid')['total_price'].sum())
```

```
df_uk['customer_monetary'] = df_uk.groupby('customerid')['total_price'].transform('sum')
```

Methodology

RFM Table

Analysis

EDA

RFM Analysis

K-Means

Cohort Analysis

```
customer_rfm = pd.merge(pd.merge(customer_recency, customer_frequency, on='customerid'), customer_monetary, on='customerid')
customer_rfm.head()
```

	customerid	recency	frequency	monetary
0	12346.000	325	1	77183.600
1	12747.000	2	11	4196.010
2	12748.000	0	209	33053.190
3	12749.000	3	5	4090.880
4	12820.000	3	4	942.340

```
customer_rfm.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 3920 entries, 0 to 3919
Data columns (total 4 columns):
#   Column      Non-Null Count  Dtype
---  -
0   customerid  3920 non-null   float64
1   recency     3920 non-null   int64
2   frequency   3920 non-null   int64
3   monetary    3920 non-null   float64
dtypes: float64(2), int64(2)
memory usage: 153.1 KB
```

Methodology

RFM Table

Analysis

EDA

RFM Analysis

K-Means

Cohort Analysis

```
quantiles = customer_rfm.quantile(q = [0.25,0.50,0.75])
quantiles
```

	recency	frequency	monetary
0.250	17.000	1.000	298.185
0.500	50.000	2.000	644.975
0.750	142.000	5.000	1571.285

	recency	frequency	monetary	recency_score	frequency_score	monetary_score
customerid						
12346.000	325	1	77183.600	1	1	4
12747.000	2	11	4196.010	4	4	4
12748.000	0	209	33053.190	4	4	4
12749.000	3	5	4090.880	4	3	4
12820.000	3	4	942.340	4	3	3

```
def RScore(x,p,d):
    if x <= d[p][0.25]:
        return 4
    elif x <= d[p][0.50]:
        return 3
    elif x <= d[p][0.75]:
        return 2
    else:
        return 1
```

```
def FMScore(x,p,d):
    if x <= d[p][0.25]:
        return 1
    elif x <= d[p][0.50]:
        return 2
    elif x <= d[p][0.75]:
        return 3
    else:
        return 4
```

Methodology

RFM Table

Analysis

EDA

RFM Analysis

K-Means

Cohort Analysis

	recency	frequency	monetary	recency_score	frequency_score	monetary_score	RFM_score	RFM_level
customerid								
14498.0	42	5	1957.32	3	3	4	334	10
17692.0	127	4	740.94	2	3	3	233	8
16209.0	88	5	2262.62	2	3	4	234	9
15277.0	46	1	255.90	3	1	1	311	5
18058.0	9	1	170.16	4	1	1	411	6
16567.0	194	2	865.60	1	2	3	123	6
13851.0	95	3	2651.46	2	3	4	234	9
14004.0	43	7	4582.64	3	4	4	344	11
15937.0	64	1	145.35	2	1	1	211	4
13294.0	189	2	873.74	1	2	3	123	6

Methodology

RFM Segmentation

Analysis

EDA

RFM Analysis

K-Means

Cohort Analysis

```
def segments(df_rfm):  
    if df_rfm['RFM_level'] == 12 :  
        return 'champion'  
    elif (df_rfm['RFM_level'] == 11) or (df_rfm['RFM_level'] == 10 ):  
        return 'loyal_customer'  
    elif (df_rfm['RFM_level'] == 9) or (df_rfm['RFM_level'] == 8 ):  
        return 'promising'  
    elif (df_rfm['RFM_level'] == 7) or (df_rfm['RFM_level'] == 6 ):  
        return 'need_attention'  
    elif (df_rfm['RFM_level'] == 5) or (df_rfm['RFM_level'] == 4 ):  
        return 'hibernating'  
    else:  
        return 'almost_lost'
```

	RFM_level	General_Segment
Bronz	6.477	776
Gold	10.462	649
Platinum	12.000	423
Silver	8.526	775
Steel	4.497	902
Trash	3.000	395

	recency	frequency	monetary	recency_score	frequency_score	monetary_score	RFM_score	RFM_level	customer_segment
customerid									
15125.0	25	15	11528.48	3	4	4	344	11	loyal_customer
12962.0	7	2	266.39	4	2	1	421	7	need_attention
16078.0	283	1	79.20	1	1	1	111	3	almost_lost
15776.0	133	1	241.62	2	1	1	211	4	hibernating
15537.0	163	1	110.92	1	1	1	111	3	almost_lost

Methodology

RFM Segmentation

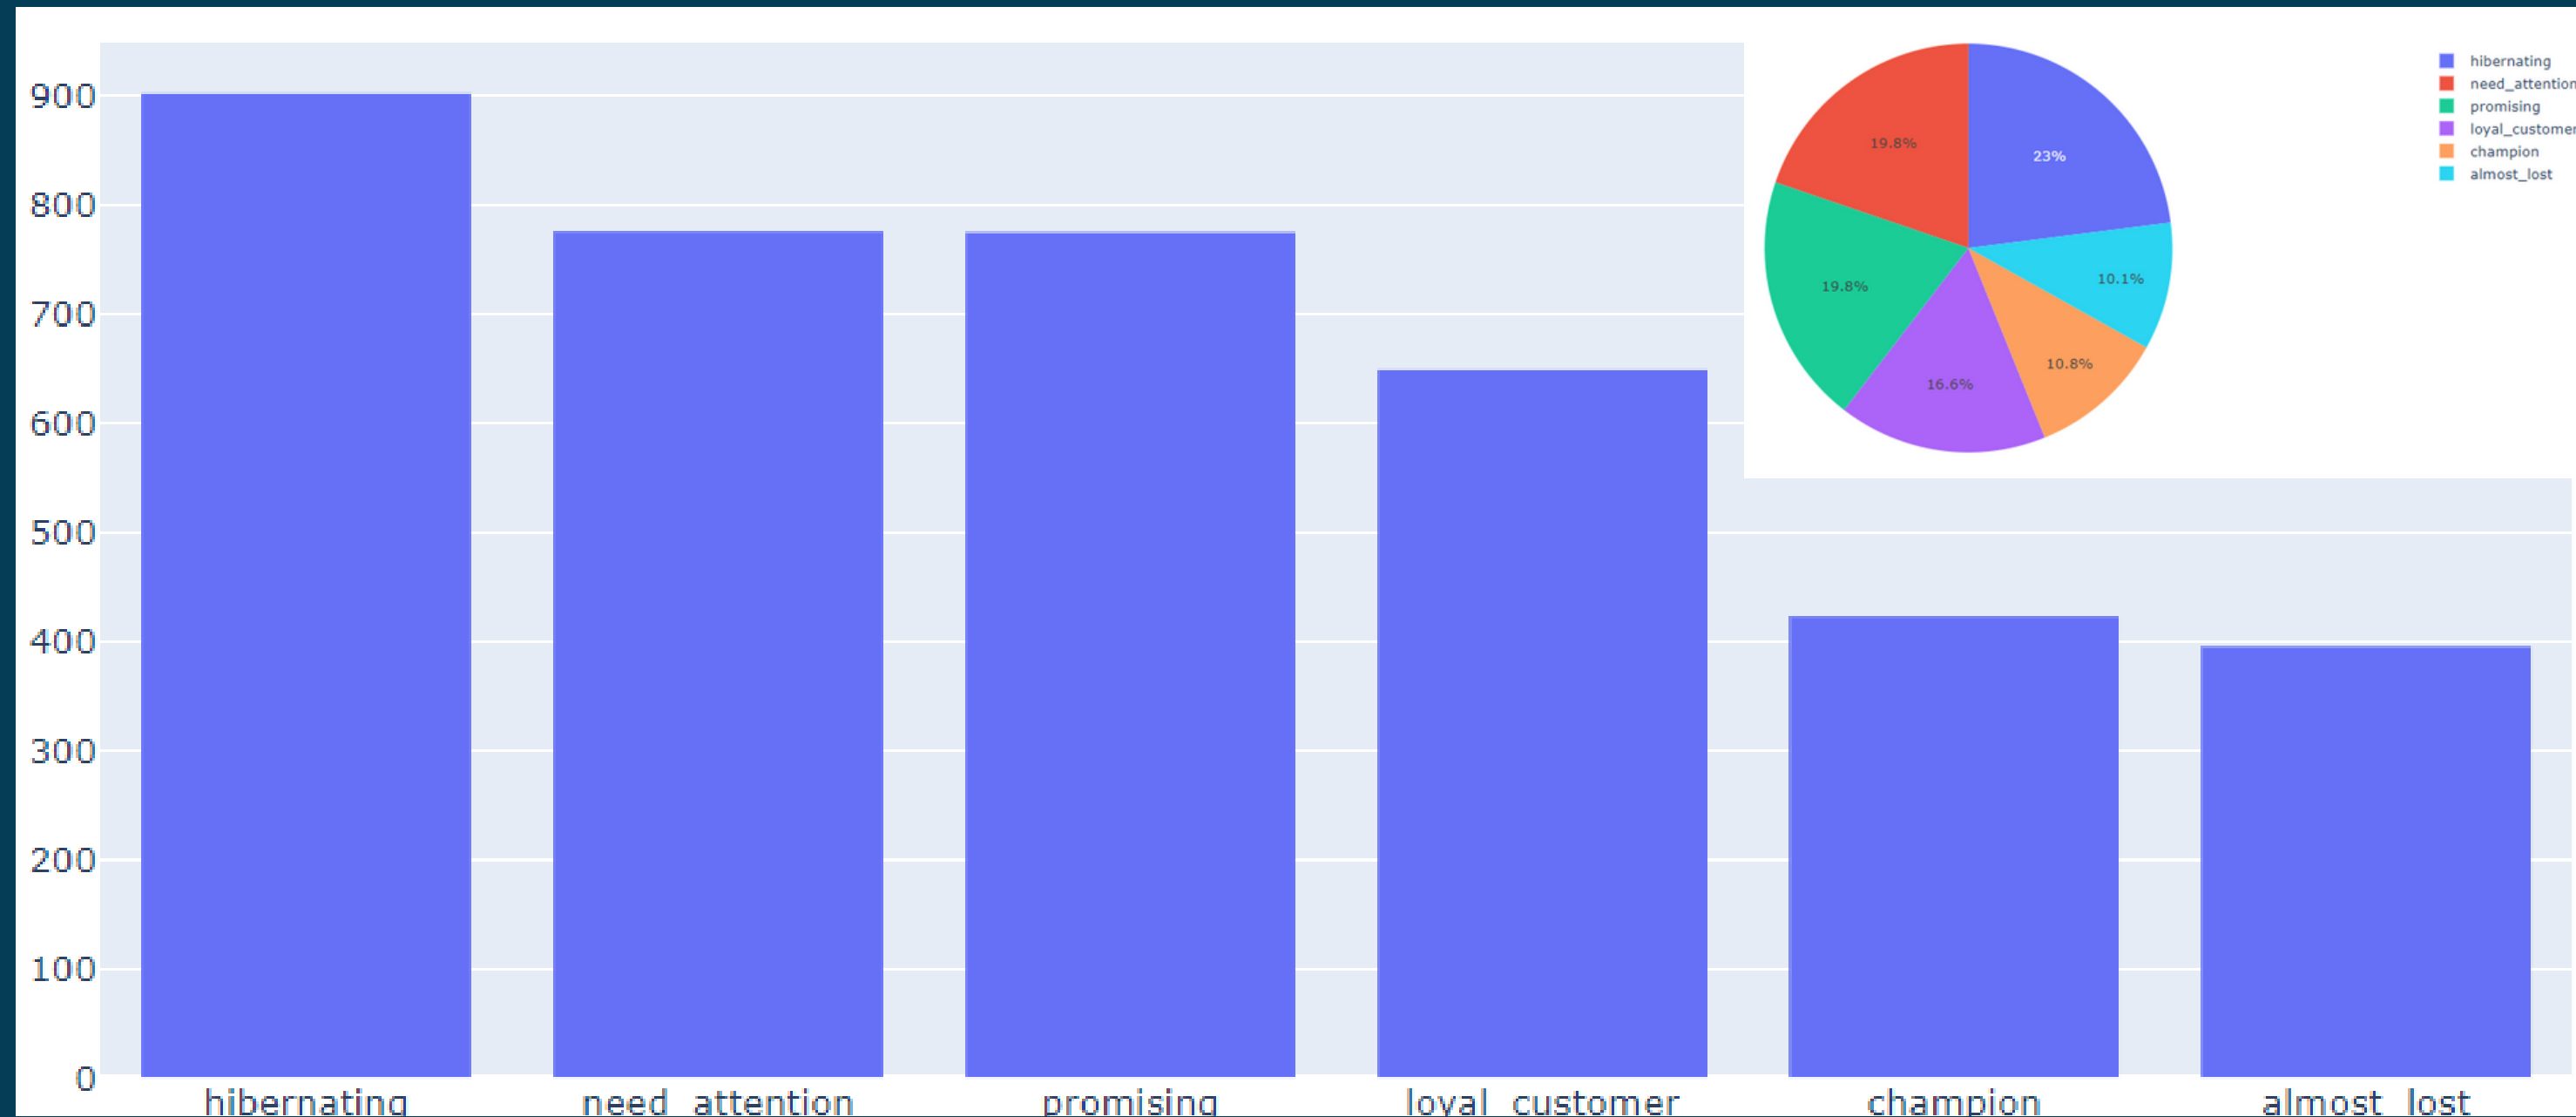
Analysis

EDA

RFM Analysis

K-Means

Cohort Analysis



Methodology

Analysis

EDA

RFM Analysis

K-Means

Cohort Analysis

```
skew_vals = customer_rfm.skew().sort_values(ascending=False)
skew_vals
```

monetary	20.217581
frequency	10.751932
recency	1.244993
frequency_score	0.153764
rfm_level	0.115653
monetary_score	0.000000
rfm_score	-0.005586
recency_score	-0.009948

```
rfm_log = customer_rfm[skew_cols.index].copy()
for col in skew_cols.index.values:
    rfm_log[col] = rfm_log[col].apply(np.log1p)
print(rfm_log[skew_cols.index].skew())
```

monetary	0.373135
frequency	1.182249
recency	-0.463737

```
rfm_before_trans = customer_rfm[skew_cols.index].copy()
pt = PowerTransformer(method='yeo-johnson')
trans = pt.fit_transform(rfm_before_trans)
rfm_after_trans = pd.DataFrame(trans, columns = skew_cols.index)
print(rfm_after_trans.skew())
```

monetary	-0.013749
frequency	0.215681
recency	-0.063910

Methodology

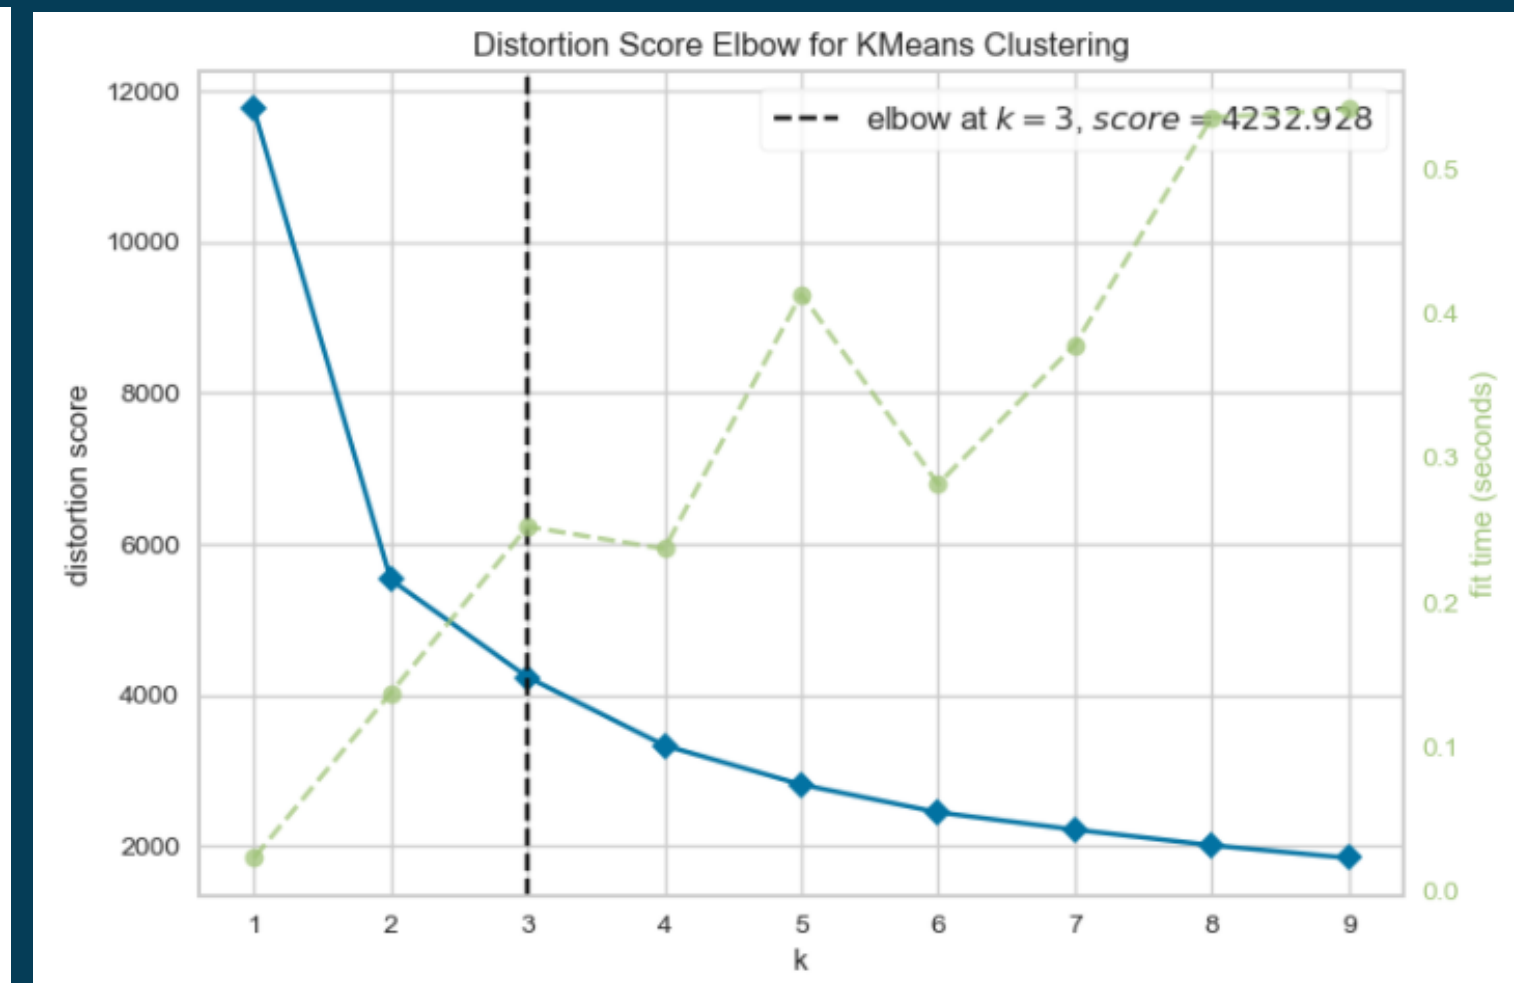

Analysis

EDA

RFM Analysis

K-Means

Cohort Analysis



```
k_means_model = KMeans(n_clusters = 3, random_state = 42)
k_means_model.fit_predict(rfm_scaled)
labels = k_means_model.labels_
rfm_trans['predicted_clusters'] = labels
rfm_trans
```

Methodology

Analysis

EDA

RFM Analysis

K-Means

Cohort Analysis

App Launched ↓		% Active users after App Launches →										
Cohort	Users	Day 0	Day 1	Day 2	Day 3	Day 4	Day 5	Day 6	Day 7	Day 8	Day 9	Day 10
Jan 25	1,098	100%	33.9%	23.5%	18.7%	15.9%	16.3%	14.2%	14.5%	Retention over user lifetime		12.1%
Jan 26	1,358	100%	31.1%	18.6%	14.3%	16.0%	14.9%	13.2%	12.9%			
Jan 27	1,257	100%	27.2%	19.6%	14.5%	12.9%	13.4%	13.0%	10.8%	11.4%		
Jan 28	1,587	100%	26.6%	17.9%	14.6%	14.8%	14.9%	13.7%	11.9%			
Jan 29	1,758	100%	26.2%	20.4%	16.9%	14.3%	12.7%	12.5%				
Jan 30	1,624	100%	26.4%	18.1%	13.7%	15.4%	11.8%					
Jan 31	1,541	100%	23.9%	19.6%	15.0%	14.8%						
Feb 01	868	100%	24.7%	16.9%	15.8%							
Feb 02	1,143	Retention over product lifetime		18.5%								
Feb 03	1,253											
All Users	13,487	100%	27.0%	19.2%	15.4%	14.9%	14.0%	13.3%	12.5%	13.1%	12.2%	12.1%

- **Acquisition Cohorts:** divide users by when they signed up first for your product. For your app users, you might break down your cohorts by the day, the week or the month they launched an app, and thereby track daily, weekly or monthly cohorts.
- **Behavioral Cohorts:** divide users by the behaviors they have (or haven't) taken in your app within a given time period. These could be any number of discrete actions that a user can perform – App Install, App Launch, App Uninstall, Transaction or Charged, or any combination of these actions / events.

Methodology

Analysis

EDA

RFM Analysis

K-Means

Cohort Analysis

ref_date	date	last_purchase_date	customer_recency	recency_value	customer_frequency	customer_monetary	invoice_month	cohort_month	cohort_index
2011-12-09	2010-12-01	2010-12-02	372 days	372	297	5391.21	2010-12-01	2010-12-01	1
2011-12-09	2010-12-01	2010-12-02	372 days	372	297	5391.21	2010-12-01	2010-12-01	1
2011-12-09	2010-12-01	2010-12-02	372 days	372	297	5391.21	2010-12-01	2010-12-01	1
2011-12-09	2010-12-01	2010-12-02	372 days	372	297	5391.21	2010-12-01	2010-12-01	1

[illegible][illegible]